

Page 1 - Executive Summary

Project Orion was completed for FinBank in 2024. The engagement was led by Sarah Chen and David Park to build an enterprise AI knowledge platform.

Page 2 - Business Context

FinBank relied on SharePoint, Confluence, and internal wikis. Knowledge discovery was identified as a strategic challenge.

Page 3 - Technology Stack

Selected technologies included Azure OpenAI, Neo4j, Databricks, and Elasticsearch. The architecture later influenced Project Atlas.

Page 4 - Knowledge Graph Design

Entities included Employee, Project, Client, Deliverable, Technology, and Regulation. Relationships included WORKED_ON, AUTHORED, IMPLEMENTED, and DEPENDS_ON.

Page 5 - Expert Discovery

The graph enabled discovery of experts. Emma Wilson was identified as an AML expert due to participation in Project Orion and the AML Modernization Initiative.

Page 6 - Document Intelligence

The platform indexed Strategy Decks, Meeting Notes, Research Reports, and Architecture Documents. Databricks pipelines generated embeddings.

Page 7 - Results

Average search time decreased by 63%. Document reuse increased by 42%. Knowledge discovery improved substantially.

Page 8 - Lessons Learned

Section-based chunking outperformed fixed-size chunking. GraphRAG retrieval outperformed vector-only retrieval. Findings were documented by David Park.

Page 9 - Future Enhancements

Recommendations included Expert Recommendation Engine, Knowledge Gap Detection, and Temporal Knowledge Graph. These later influenced Project Atlas.

Page 10 - Conclusion

Project Orion demonstrated that enterprise knowledge graphs improve institutional memory. Sarah Chen recommended expansion across additional business units.